

✓ Mount Google Drive and Set Path

```
from google.colab import drive
drive.mount('/content/drive')
```

⇄ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount(

```
import os
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from collections import Counter

# Define path
base_dir = '/content/drive/MyDrive/Assessment/Assessment_1/dataset'
train_dir = os.path.join(base_dir, 'train')
test_dir = os.path.join(base_dir, 'test')
```

✓ Set Parameters and Create Data Generators

```
IMG_SIZE = (224, 224)
BATCH_SIZE = 32

# No augmentation
train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

# Load images
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode='categorical'
)

test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode='categorical',
    shuffle=False
)
```

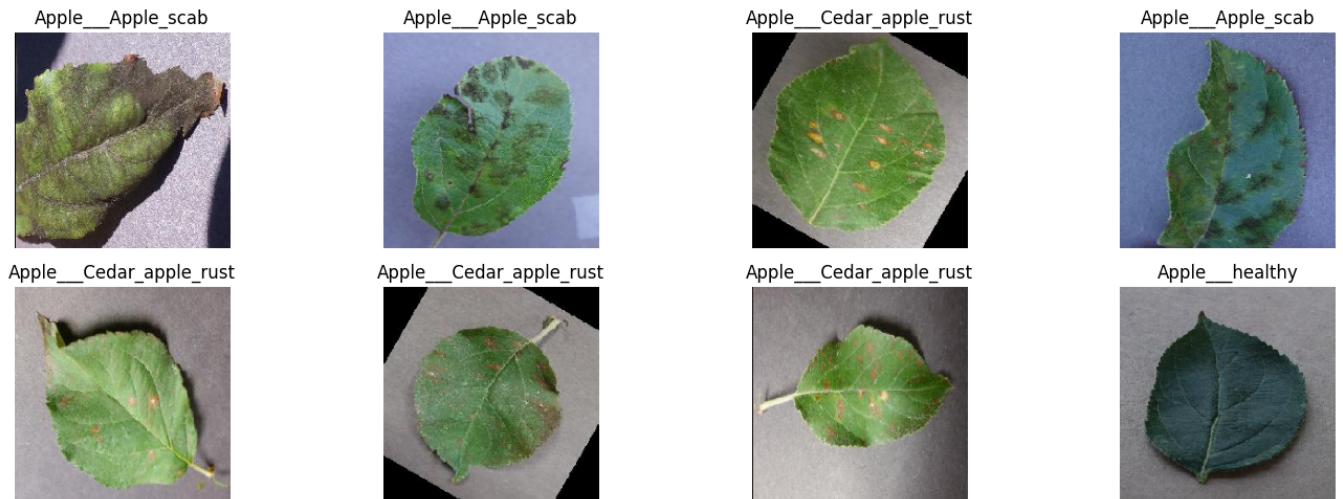


```
Found 6218 images belonging to 4 classes.
Found 1553 images belonging to 4 classes.
```


✓ Visualize Sample Images

```
def plot_images(generator, class_indices, num_images=8):
    class_names = list(class_indices.keys())
    x_batch, y_batch = next(generator)
    plt.figure(figsize=(15, 5))
    for i in range(num_images):
        ax = plt.subplot(2, 4, i + 1)
        plt.imshow(x_batch[i])
        label_index = np.argmax(y_batch[i])
        plt.title(class_names[label_index])
        plt.axis("off")
    plt.tight_layout()
    plt.show()


plot_images(train_generator, train_generator.class_indices)
```



Dataset Analysis



Test Dataset Class Distribution



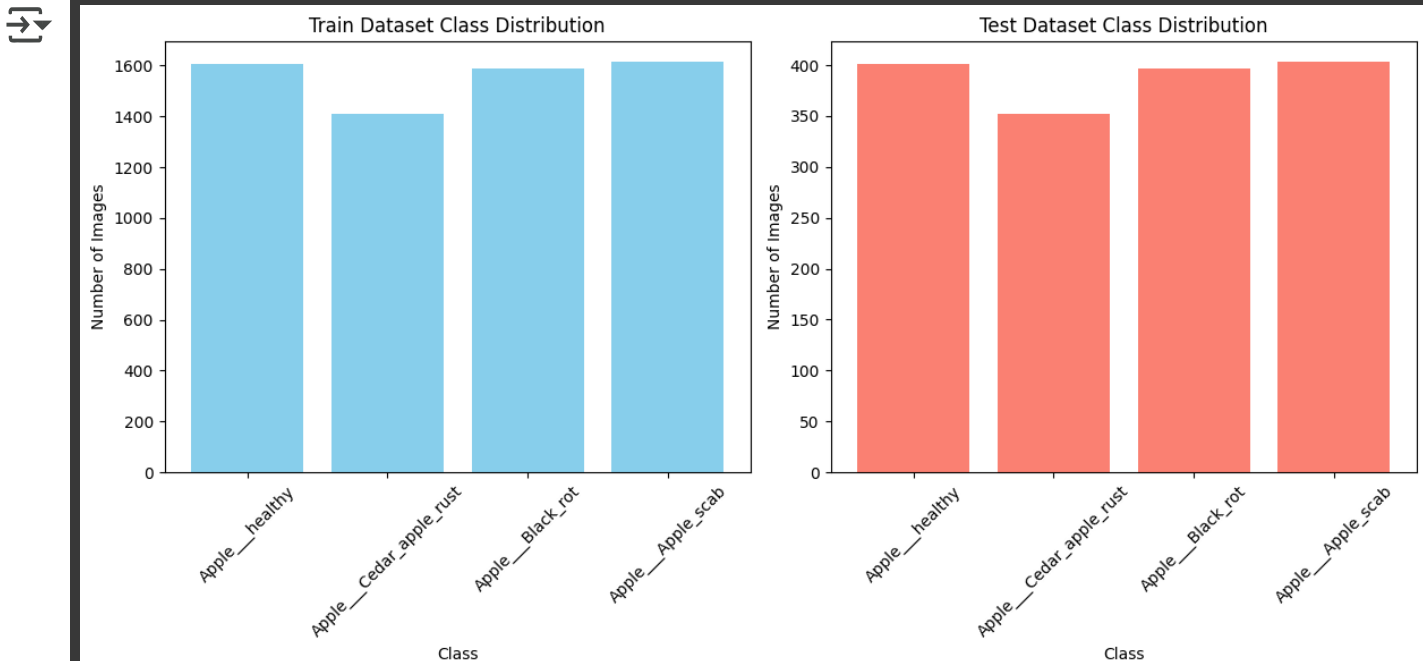
✓ Bar Plot for Class-wise Image Counts

```
# Plot class distribution using bar charts
plt.figure(figsize=(12, 6))

# Train dataset bar plot
plt.subplot(1, 2, 1)
plt.bar(train_counts.keys(), train_class_counts, color='skyblue')
plt.title("Train Dataset Class Distribution")
plt.xlabel("Class")
plt.ylabel("Number of Images")
plt.xticks(rotation=45)

# Test dataset bar plot
plt.subplot(1, 2, 2)
plt.bar(test_counts.keys(), test_class_counts, color='salmon')
plt.title("Test Dataset Class Distribution")
plt.xlabel("Class")
plt.ylabel("Number of Images")
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```



Filter A filter (or feature detector) refers to the set of values (weights) that are applied to the input data during the convolution process.

Kernel A kernel is the actual matrix (the set of weights) that is used to perform the convolution operation on the input image. It is typically smaller than the image (e.g., a 3x3 or 5x5 matrix)

✓ Build the Baseline CNN Architecture

Convolutional Layer: Extracts local features using filters. Activation Layer (ReLU): Introduces non-linearity. Pooling Layer: Reduces spatial dimensions and retains essential features. Flatten Layer: Converts 2D feature maps into 1D vectors. Fully Connected Layer: Learns high-level features for classification or regression. Dropout Layer: Prevents overfitting by randomly deactivating neurons. Batch Normalization: Normalizes layer inputs for more stable training.

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import Adam

# Build the model
model = Sequential() #This initializes a Sequential model

# Convolutional Layer 1
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3))) #32 is number of f:
model.add(MaxPooling2D(pool_size=(2, 2)))

# Convolutional Layer 2
model.add(Conv2D(64, (3, 3), activation='relu')) #(3, 3): The size of each filter (kernel)
model.add(MaxPooling2D(pool_size=(2, 2)))

# Convolutional Layer 3
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

# Flatten layer
model.add(Flatten())
#flatten the multi-dimensional output from the previous layer (typically the 2D feature maps a
# Fully connected layers
model.add(Dense(512, activation='relu')) #learns the feature
model.add(Dropout(0.5)) # Dropout to prevent overfitting
model.add(Dense(256, activation='relu'))
model.add(Dense(4, activation='softmax')) # 4 classes

# Model summary
model.summary()

```



Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d_3 (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_4 (Conv2D)	(None, 109, 109, 64)	18,496
max_pooling2d_4 (MaxPooling2D)	(None, 54, 54, 64)	0
conv2d_5 (Conv2D)	(None, 52, 52, 128)	73,856
max_pooling2d_5 (MaxPooling2D)	(None, 26, 26, 128)	0
flatten_1 (Flatten)	(None, 86528)	0
dense_3 (Dense)	(None, 512)	44,302,848
dropout_1 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 256)	131,328
dense_5 (Dense)	(None, 4)	1,028

Total params: 44,528,452 (169.86 MB)
 Trainable params: 44,528,452 (169.86 MB)
 Non-trainable params: 0 (0.00 B)

A Sequential model is a linear stack of layers where each layer has exactly one input and one output

✓ Compile the Model


```
# Compile the model
model.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
```

✓ Train the Model

```
history = model.fit(
    train_generator,
    epochs=10, # Adjust as necessary
    validation_data=test_generator,
    verbose=1 #This controls how much information is shown during training.
)

# Plotting training & validation loss
plt.figure(figsize=(8, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training vs Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

```
→ /usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapte
    self._warn_if_super_not_called()
Epoch 1/10
195/195 ————— 4867s 25s/step - accuracy: 0.6507 - loss: 0.8926 - val_accura
Epoch 2/10
195/195 ————— 37s 190ms/step - accuracy: 0.8997 - loss: 0.2961 - val_accura
Epoch 3/10
195/195 ————— 34s 174ms/step - accuracy: 0.9277 - loss: 0.1924 - val_accura
Epoch 4/10
195/195 ————— 40s 168ms/step - accuracy: 0.9505 - loss: 0.1561 - val_accura
Epoch 5/10
195/195 ————— 42s 174ms/step - accuracy: 0.9631 - loss: 0.1060 - val_accura
Epoch 6/10
195/195 ————— 33s 168ms/step - accuracy: 0.9684 - loss: 0.0886 - val_accura
Epoch 7/10
195/195 ————— 37s 188ms/step - accuracy: 0.9804 - loss: 0.0586 - val_accura
Epoch 8/10
195/195 ————— 32s 165ms/step - accuracy: 0.9835 - loss: 0.0479 - val_accura
Epoch 9/10
195/195 ————— 33s 169ms/step - accuracy: 0.9772 - loss: 0.0661 - val_accura
Epoch 10/10
195/195 ————— 32s 165ms/step - accuracy: 0.9821 - loss: 0.0420 - val_accura
```



✓ Evaluate the Model

```

from sklearn.metrics import classification_report

# Get predictions
predictions = model.predict(test_generator)
y_pred = np.argmax(predictions, axis=1) #picks the class with the highest probability for each
y_true = test_generator.classes

# Print classification report
print(classification_report(y_true, y_pred, target_names=test_generator.class_indices.keys()))

```

↔ 49/49 ————— 7s 141ms/step

	precision	recall	f1-score	support
Apple___Apple_scab	0.98	0.90	0.94	403
Apple___Black_rot	0.93	0.99	0.96	397
Apple___Cedar_apple_rust	0.97	1.00	0.98	352
Apple___healthy	0.98	0.97	0.97	401
accuracy			0.96	1553
macro avg	0.96	0.97	0.96	1553
weighted avg	0.96	0.96	0.96	1553

Precision measures the accuracy of positive predictions. Recall measures how well the model identifies actual positive instances. F1-Score is the harmonic mean of precision and recall.

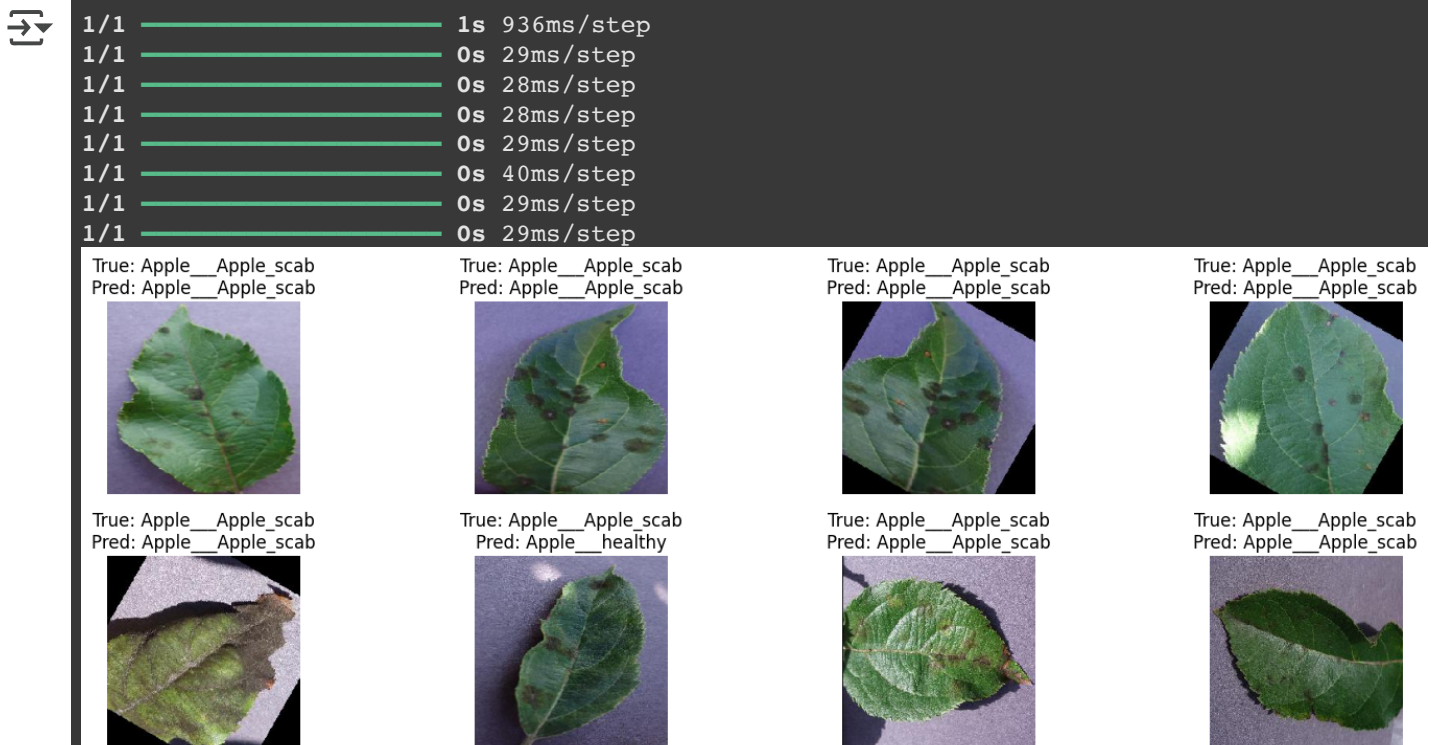
✓ Visualize Predictions

```
def plot_predictions(generator, model, num_images=8):
    class_names = list(generator.class_indices.keys())
    x_batch, y_batch = next(generator)
    plt.figure(figsize=(15, 5))

    for i in range(num_images):
        ax = plt.subplot(2, 4, i + 1)
        plt.imshow(x_batch[i])
        predicted_class = class_names[np.argmax(model.predict(x_batch[i:i+1]))]
        true_class = class_names[np.argmax(y_batch[i])]
        plt.title(f"True: {true_class}\nPred: {predicted_class}")
        plt.axis('off')

    plt.tight_layout()
    plt.show()

plot_predictions(test_generator, model)
```



batch normalization is normalizes the inputs to each layer of the network. This means it adjusts the data so that each input has a mean of 0 and a variance of 1. This makes the training process smoother and faster.

✓ Build the Deeper Model

```
from tensorflow.keras.layers import BatchNormalization

# Build deeper model
model_deeper = Sequential()
```

```

# Convolutional Layer 1
model_deeper.add(Conv2D(64, (3, 3), activation='relu', input_shape=(224, 224, 3)))
model_deeper.add(BatchNormalization()) # Batch normalization
model_deeper.add(MaxPooling2D(pool_size=(2, 2)))

# Convolutional Layer 2
model_deeper.add(Conv2D(128, (3, 3), activation='relu'))
model_deeper.add(BatchNormalization())
model_deeper.add(MaxPooling2D(pool_size=(2, 2)))

# Convolutional Layer 3
model_deeper.add(Conv2D(256, (3, 3), activation='relu'))
model_deeper.add(BatchNormalization())
model_deeper.add(MaxPooling2D(pool_size=(2, 2)))


# Convolutional Layer 4 (Added layer)
model_deeper.add(Conv2D(512, (3, 3), activation='relu'))
model_deeper.add(BatchNormalization())
model_deeper.add(MaxPooling2D(pool_size=(2, 2)))

# Convolutional Layer 5 (Added layer)
model_deeper.add(Conv2D(512, (3, 3), activation='relu'))
model_deeper.add(BatchNormalization())
model_deeper.add(MaxPooling2D(pool_size=(2, 2)))

# Flatten layer
model_deeper.add(Flatten())

# Fully connected layers
model_deeper.add(Dense(1024, activation='relu'))
model_deeper.add(Dropout(0.5)) # Dropout for regularization
model_deeper.add(Dense(512, activation='relu'))
model_deeper.add(Dense(4, activation='softmax')) # 4 classes
#decreasing will neurons to 64, you are decreasing the model's capacity.
# Model summary
model_deeper.summary()

```

 /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: U
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 222, 222, 64)	1,792
batch_normalization (BatchNormalization)	(None, 222, 222, 64)	256
max_pooling2d_6 (MaxPooling2D)	(None, 111, 111, 64)	0
conv2d_7 (Conv2D)	(None, 109, 109, 128)	73,856
batch_normalization_1 (BatchNormalization)	(None, 109, 109, 128)	512
max_pooling2d_7 (MaxPooling2D)	(None, 54, 54, 128)	0
conv2d_8 (Conv2D)	(None, 52, 52, 256)	295,168
batch_normalization_2 (BatchNormalization)	(None, 52, 52, 256)	1,024
max_pooling2d_8 (MaxPooling2D)	(None, 26, 26, 256)	0
conv2d_9 (Conv2D)	(None, 24, 24, 512)	1,180,160
batch_normalization_3 (BatchNormalization)	(None, 24, 24, 512)	2,048
max_pooling2d_9 (MaxPooling2D)	(None, 12, 12, 512)	0
conv2d_10 (Conv2D)	(None, 10, 10, 512)	2,359,808
batch_normalization_4 (BatchNormalization)	(None, 10, 10, 512)	2,048
max_pooling2d_10 (MaxPooling2D)	(None, 5, 5, 512)	0
flatten_2 (Flatten)	(None, 12800)	0
dense_6 (Dense)	(None, 1024)	13,108,224
dropout_2 (Dropout)	(None, 1024)	0
dense_7 (Dense)	(None, 512)	524,800
dense_8 (Dense)	(None, 4)	2,052

Total params: 17,551,748 (66.95 MB)

Trainable params: 17,548,804 (66.94 MB)

Non-trainable params: 2,944 (11.50 KB)

✓ Compile the Model

```
# Compile the deeper model
model_deeper.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
```


✓ Train the Deeper Model

```
history_deeper = model_deeper.fit(
    train_generator,
    epochs=10, # Adjust as necessary
    validation_data=test_generator,
    verbose=1 #This controls how much information is shown during training
)

# Plotting training & validation loss for deeper model
plt.figure(figsize=(8, 6))
plt.plot(history_deeper.history['loss'], label='Training Loss')
plt.plot(history_deeper.history['val_loss'], label='Validation Loss')
plt.title('Training vs Validation Loss (Deeper Model)')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
Epoch 1/10
195/195 ————— 69s 263ms/step - accuracy: 0.6703 - loss: 3.8081 - val_accu
Epoch 2/10
195/195 ————— 36s 186ms/step - accuracy: 0.8610 - loss: 0.4941 - val_accu
Epoch 3/10
195/195 ————— 38s 194ms/step - accuracy: 0.9216 - loss: 0.2724 - val_accu
Epoch 4/10
195/195 ————— 36s 185ms/step - accuracy: 0.9310 - loss: 0.2050 - val_accu
Epoch 5/10
195/195 ————— 36s 184ms/step - accuracy: 0.9500 - loss: 0.1476 - val_accu
Epoch 6/10
195/195 ————— 36s 181ms/step - accuracy: 0.9663 - loss: 0.1046 - val_accu
Epoch 7/10
195/195 ————— 36s 185ms/step - accuracy: 0.9709 - loss: 0.0861 - val_accu
Epoch 8/10
195/195 ————— 36s 182ms/step - accuracy: 0.9645 - loss: 0.1154 - val_accu
Epoch 9/10
195/195 ————— 35s 181ms/step - accuracy: 0.9729 - loss: 0.0962 - val_accu
Epoch 10/10
195/195 ————— 39s 201ms/step - accuracy: 0.9694 - loss: 0.0955 - val_accu
```



✓ Evaluate the Deeper Model


```
# Get predictions from deeper model
predictions_deeper = model_deeper.predict(test_generator)
y_pred_deeper = np.argmax(predictions_deeper, axis=1)

# Print classification report
from sklearn.metrics import classification_report

print(classification_report(y_true, y_pred_deeper, target_names=test_generator.class_indices.keys()))
```

```
⇒ 49/49 ————— 8s 144ms/step
              precision    recall  f1-score   support









 Apple__Apple_scab         0.75      0.95      0.84         403
 Apple__Black_rot          0.96      0.99      0.98         397
 Apple__Cedar_apple_rust    0.84      0.82      0.83         352
 Apple__healthy            0.97      0.70      0.81         401

      accuracy              0.87         1553
      macro avg              0.88         1553
      weighted avg           0.88         1553
```

✓ Visualize Predictions from the Deeper Model

```
plot_predictions(test_generator, model_deeper)
```

```
⇒ 1/1 ————— 1s 1s/step
   1/1 ————— 0s 30ms/step
   1/1 ————— 0s 29ms/step
   1/1 ————— 0s 29ms/step
   1/1 ————— 0s 33ms/step
   1/1 ————— 0s 30ms/step
   1/1 ————— 0s 30ms/step
   1/1 ————— 0s 32ms/step
```

<p>True: Apple__Apple_scab Pred: Apple__Apple_scab</p> 	<p>True: Apple__Apple_scab Pred: Apple__Apple_scab</p> 	<p>True: Apple__Apple_scab Pred: Apple__Apple_scab</p> 	<p>True: Apple__Apple_scab Pred: Apple__Apple_scab</p> 
<p>True: Apple__Apple_scab Pred: Apple__Apple_scab</p> 	<p>True: Apple__Apple_scab Pred: Apple__Apple_scab</p> 	<p>True: Apple__Apple_scab Pred: Apple__Apple_scab</p> 	<p>True: Apple__Apple_scab Pred: Apple__Apple_scab</p> 


Comparative Analysis of Baseline vs. Deeper Model


✓ Model Performance Comparison

```
# Compare accuracy and loss for baseline and deeper models
baseline_history = history.history
deeper_history = history_deeper.history


# Plot accuracy comparison
plt.figure(figsize=(8, 6))
plt.plot(baseline_history['accuracy'], label='Baseline Training Accuracy')
plt.plot(deeper_history['accuracy'], label='Deeper Training Accuracy')
plt.plot(baseline_history['val_accuracy'], label='Baseline Validation Accuracy')
plt.plot(deeper_history['val_accuracy'], label='Deeper Validation Accuracy')
plt.title('Training and Validation Accuracy Comparison')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Plot loss comparison
plt.figure(figsize=(8, 6))
plt.plot(baseline_history['loss'], label='Baseline Training Loss')
plt.plot(deeper_history['loss'], label='Deeper Training Loss')
plt.plot(baseline_history['val_loss'], label='Baseline Validation Loss')
plt.plot(deeper_history['val_loss'], label='Deeper Validation Loss')
plt.title('Training and Validation Loss Comparison')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```





Training and Validation Loss Comparison



Computational Efficiency Comparison

✓ Optimizer Comparison (SGD vs Adam)

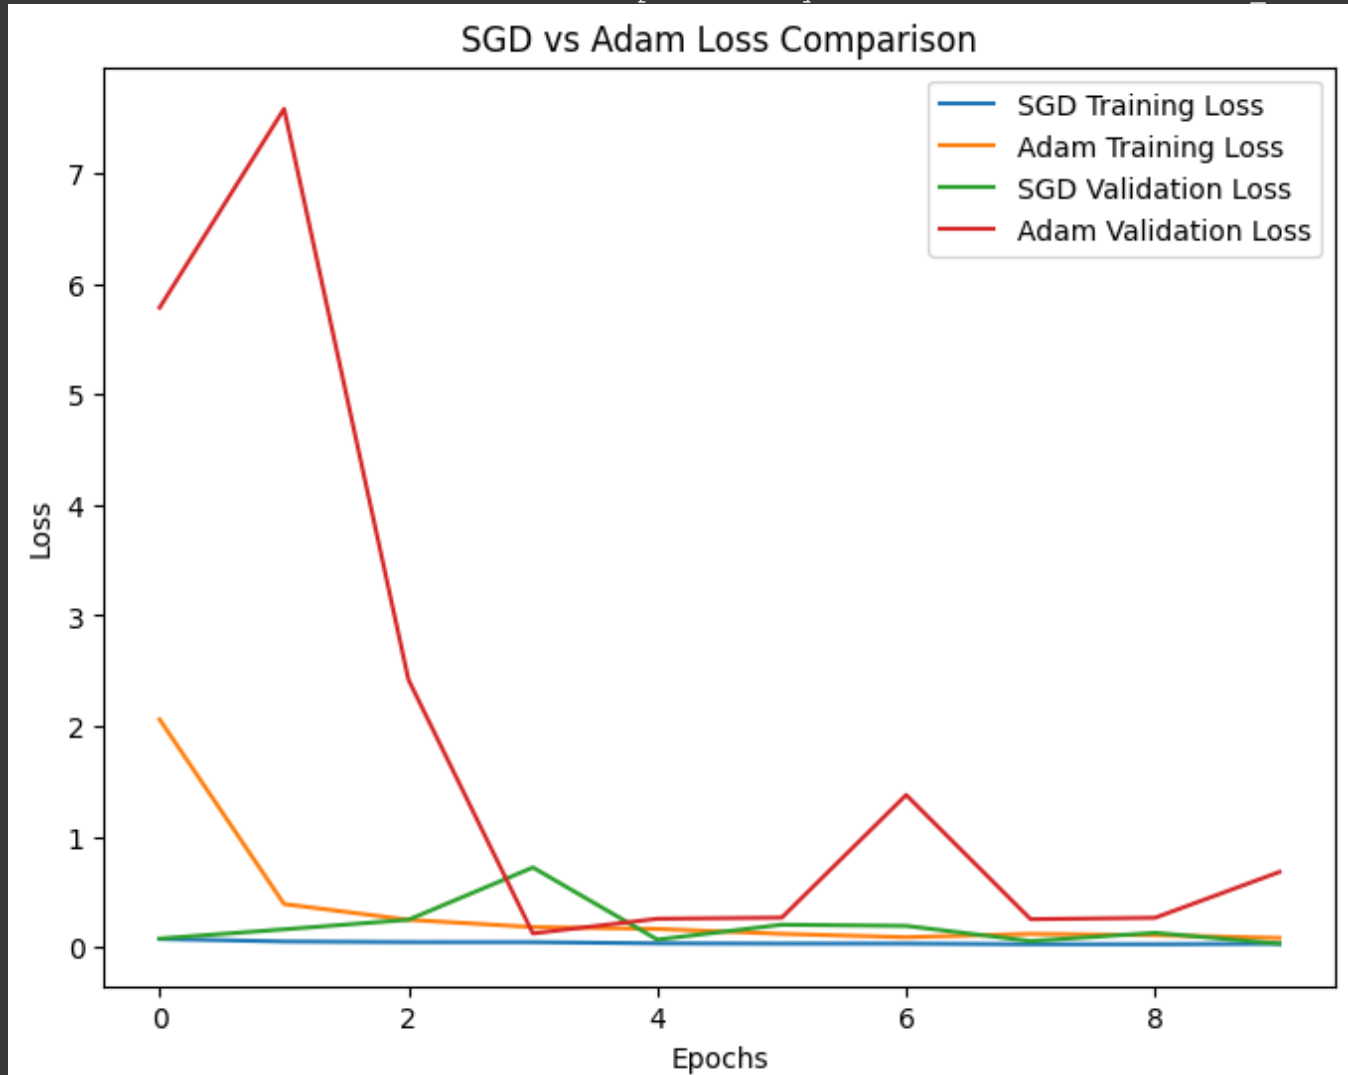
```
# Compile model with SGD optimizer
model_deeper.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Train with SGD learning rate is 0.01 meanwhile adam is 0.0005
history_sgd = model_deeper.fit(
    train_generator,
    epochs=10, #one complete pass through the entire training dataset during the training process
    validation_data=test_generator,
    verbose=1
)

# Plot SGD vs Adam Loss
plt.figure(figsize=(8, 6))
plt.plot(history_sgd.history['loss'], label='SGD Training Loss')
```

```
plt.plot(deeper_history['loss'], label='Adam Training Loss')
plt.plot(history_sgd.history['val_loss'], label='SGD Validation Loss')
plt.plot(deeper_history['val_loss'], label='Adam Validation Loss')
plt.title('SGD vs Adam Loss Comparison')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

Epoch 1/10
195/195 ————— 43s 199ms/step - accuracy: 0.9754 - loss: 0.0951 - val_accu
Epoch 2/10
195/195 ————— 35s 180ms/step - accuracy: 0.9844 - loss: 0.0488 - val_accu
Epoch 3/10
195/195 ————— 42s 183ms/step - accuracy: 0.9873 - loss: 0.0430 - val_accu
Epoch 4/10
195/195 ————— 35s 180ms/step - accuracy: 0.9890 - loss: 0.0410 - val_accu
Epoch 5/10
195/195 ————— 36s 181ms/step - accuracy: 0.9895 - loss: 0.0300 - val_accu
Epoch 6/10
195/195 ————— 35s 180ms/step - accuracy: 0.9911 - loss: 0.0282 - val_accu
Epoch 7/10
195/195 ————— 44s 198ms/step - accuracy: 0.9875 - loss: 0.0314 - val_accu
Epoch 8/10
195/195 ————— 36s 184ms/step - accuracy: 0.9896 - loss: 0.0248 - val_accu
Epoch 9/10
195/195 ————— 35s 178ms/step - accuracy: 0.9936 - loss: 0.0252 - val_accu
Epoch 10/10
195/195 ————— 35s 180ms/step - accuracy: 0.9881 - loss: 0.0334 - val_accu



✓ Loading and Adapting a Pre-Trained Model

```
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import GlobalAveragePooling2D
```


```
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam

# Load VGG16 with pre-trained ImageNet weights, without the top fully connected layers
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
#will only retain the convolutional layers and pooling layers, which are responsible for feature extraction
# Freeze the base model (so the weights won't be updated during training)
for layer in base_model.layers:
    layer.trainable = False

# Add a custom top layer for your classification task
x = base_model.output
x = GlobalAveragePooling2D()(x) # Pooling layer to reduce dimensionality
x = Dense(1024, activation='relu')(x)
x = Dropout(0.5)(x) # Regularization
x = Dense(4, activation='softmax')(x) # 4 classes for classification

# Create the final model
model_transfer = Model(inputs=base_model.input, outputs=x)

# Compile the model
model_transfer.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model_transfer.summary()
#Adam(learning_rate=0.0005)
```



Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/vgg16/v58889256/58889256> 4s 0us/step
 Model: "functional_42"

Layer (type)	Output Shape	Param #
input_layer_3 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense_9 (Dense)	(None, 1024)	525,312
dropout_3 (Dropout)	(None, 1024)	0
dense_10 (Dense)	(None, 4)	4,100

Total params: 15,244,100 (58.15 MB)
 Trainable params: 529,412 (2.02 MB)
 Non-trainable params: 14,714,688 (56.13 MB)

Training Strategies

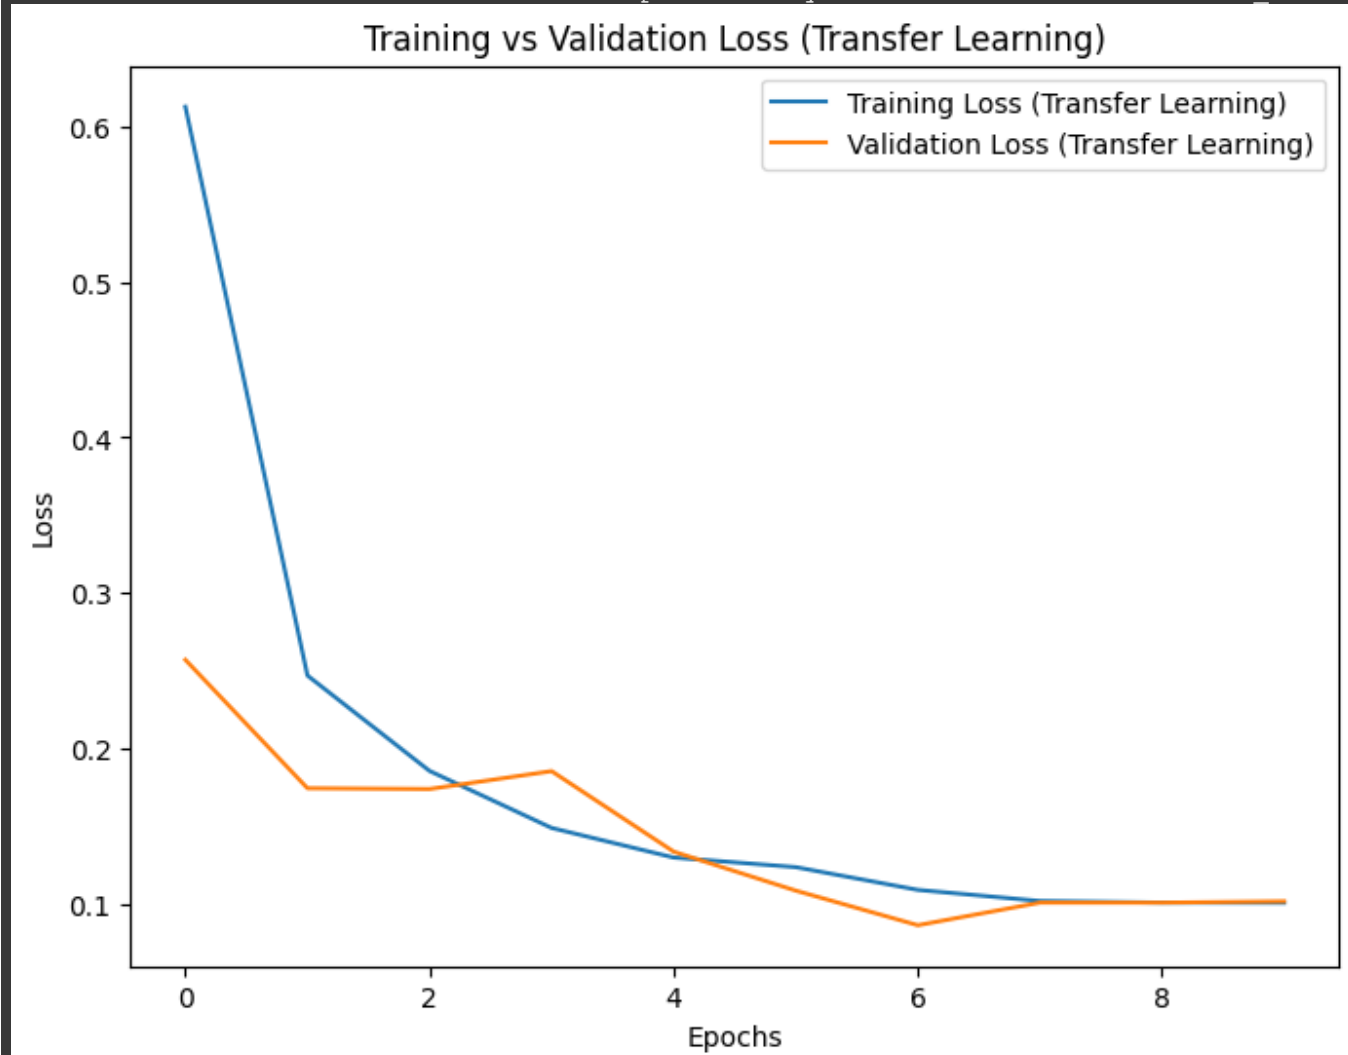
```

# Train only the custom top layers
history_transfer = model_transfer.fit(
    train_generator,
    epochs=10, # You can adjust this based on training time
  
```

```
        validation_data=test_generator,  
        verbose=1 #This controls the level of logging output during training  
    )  
  
# Plot training vs validation loss for transfer learning  
plt.figure(figsize=(8, 6))  
plt.plot(history_transfer.history['loss'], label='Training Loss (Transfer Learning)')  
plt.plot(history_transfer.history['val_loss'], label='Validation Loss (Transfer Learning)')  
plt.title('Training vs Validation Loss (Transfer Learning)')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```




```
Epoch 1/10
195/195 ————— 70s 287ms/step - accuracy: 0.6281 - loss: 0.9107 - val_accu
Epoch 2/10
195/195 ————— 55s 212ms/step - accuracy: 0.9059 - loss: 0.2711 - val_accu
Epoch 3/10
195/195 ————— 82s 210ms/step - accuracy: 0.9387 - loss: 0.1827 - val_accu
Epoch 4/10
195/195 ————— 40s 206ms/step - accuracy: 0.9494 - loss: 0.1472 - val_accu
Epoch 5/10
195/195 ————— 40s 206ms/step - accuracy: 0.9425 - loss: 0.1556 - val_accu
Epoch 6/10
195/195 ————— 40s 206ms/step - accuracy: 0.9548 - loss: 0.1327 - val_accu
Epoch 7/10
195/195 ————— 40s 205ms/step - accuracy: 0.9632 - loss: 0.1091 - val_accu
Epoch 8/10
195/195 ————— 41s 206ms/step - accuracy: 0.9691 - loss: 0.0890 - val_accu
Epoch 9/10
195/195 ————— 40s 205ms/step - accuracy: 0.9610 - loss: 0.1103 - val_accu
Epoch 10/10
195/195 ————— 40s 206ms/step - accuracy: 0.9621 - loss: 0.1058 - val_accu
```



✓ Model Evaluation and Prediction

```
# Get predictions from the transfer learning model
predictions_transfer = model_transfer.predict(test_generator)
y_pred_transfer = np.argmax(predictions_transfer, axis=1)

# Print classification report for transfer learning model
print(classification_report(y_true, y_pred_transfer, target_names=test_generator.class_indices
```

```
⇒ 49/49 ————— 10s 184ms/step
```

	precision	recall	f1-score	support
Apple__Apple_scab	0.99	0.87	0.92	403
Apple__Black_rot	0.97	0.99	0.98	397
Apple__Cedar_apple_rust	0.93	0.99	0.96	352
Apple__healthy	0.94	0.99	0.96	401
accuracy			0.96	1553
macro avg	0.96	0.96	0.96	1553
weighted avg	0.96	0.96	0.96	1553

✓ ROC curve (Receiver Operating Characteristic curve)

```
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize

# Binarize the true labels (multi-class one-hot encoding)
y_true_binarized = label_binarize(y_true, classes=[0, 1, 2, 3])

# Function to plot the ROC curve for each class
def plot_roc_curve(y_true, y_pred_prob, model_name):
    n_classes = y_true.shape[1] # Number of classes
    fpr, tpr, auc_scores = {}, {}, {}

    # Plot ROC curve for each class
    plt.figure(figsize=(10, 8))

    for i in range(n_classes):
        fpr[i], tpr[i], _ = roc_curve(y_true[:, i], y_pred_prob[:, i]) # Get FPR, TPR for class i
        auc_scores[i] = auc(fpr[i], tpr[i]) # Calculate AUC for each class
        plt.plot(fpr[i], tpr[i], label=f'Class {i} (AUC = {auc_scores[i]:.2f})')

    # Plot a random classifier (diagonal line)
    plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier')

    plt.title(f'ROC Curve for {model_name}')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.legend(loc='best')
    plt.show()


# Get predicted probabilities for each model (probabilities for each class)
y_pred_baseline_prob = model.predict(test_generator)
y_pred_deeper_prob = model_deeper.predict(test_generator)
y_pred_transfer_prob = model_transfer.predict(test_generator)

# Plot ROC curve for each model
plot_roc_curve(y_true_binarized, y_pred_baseline_prob, 'Baseline Model')
plot_roc_curve(y_true_binarized, y_pred_deeper_prob, 'Deeper Model')
plot_roc_curve(y_true_binarized, y_pred_transfer_prob, 'Transfer Learning Model')
```




49/49 6s 126ms/step
49/49 8s 146ms/step
49/49 9s 173ms/step

ROC Curve for Baseline Model



ROC Curve for Deeper Model



The ROC curve plots the True Positive Rate (TPR) (on the Y-axis) against the False Positive Rate (FPR) (on the X-axis) for different thresholds, that quantifies the overall ability of the model to distinguish between positive and negative classes.

