

✓ Worksheet 8

```
from google.colab import drive
drive.mount('/content/drive')
```

⇨ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount('/content/drive', force_remount=True).

✓ Helper Function for Text Cleaning:

Implement a Helper Function as per Text Preprocessing Notebook and Complete the following pipeline.

```
import re
def remove_urls(text):
    """
    This function will try to remove URL present in our dataset and replace it with
    Input Args:
    text: strings of text that may contain URLs.
    Output Args:
    text: URLs replaces with text
    """
    url_pattern = re.compile(r'https?://\S+|www\.\S+')
    return url_pattern.sub(r'', text)
```

```
def remove_emoji(string):
    """
    This function will replace the emoji in string with whitespace
    """
    emoji_pattern = re.compile("[
        u"\U0001F600-\U0001F64F" # emoticons
        u"\U0001F300-\U0001F5FF" # symbols & pictographs
        u"\U0001F680-\U0001F6FF" # transport & map symbols
        u"\U0001F1E0-\U0001F1FF" # flags (iOS)
        u"\U00002702-\U000027B0"
        u"\U000024C2-\U0001F251"
        "]" +, flags=re.UNICODE)
    return emoji_pattern.sub(r' ', string)
```

```
def removeunwanted_characters(document):
    """
    This function will remove all the unwanted characters from the input dataset.
    Input Args:
    document: A text data to be cleaned.
    Return:
    A cleaned document.
    """
    # remove user mentions
    document = re.sub("@[A-Za-z0-9_]+", " ", document)
    # remove hashtags
    document = re.sub("#[A-Za-z0-9_]+", "", document)
    # remove punctuation
    document = re.sub("[^0-9A-Za-z ]", "" , document)
    #remove emojis
    document = remove_emoji(document)
    # remove double spaces
    document = document.replace('  ', "")
    return document.strip()
```

```
import nltk
nltk.download('punkt_tab')
from nltk import word_tokenize
```

🔄 [nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!

```
from nltk.tokenize import RegexpTokenizer

from nltk.tokenize import RegexpTokenizer

def remove_punct(text):
    """
    This function removes the punctuations present in our text data.
    Input Args:
    text: text data.
    Returns:
    text: cleaned text.
    """
    tokenizer = RegexpTokenizer(r"\w+")
    lst=tokenizer.tokenize(' '.join(text))
    return lst
```

```
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
stop_words = set(stopwords.words('english'))
custom_stopwords = ['@', 'RT']
stop_words.update(custom_stopwords)
```

```
➡ [nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
def remove_stopwords(text_tokens):
    """
    This function removes all the stopwords present in our text tokens.
    Input Args:
    text_tokens: tokenize input of our datasets.
    Returns:
    result_tokens: list of token without stopword.
    """

    result_tokens = []
    for token in text_tokens:
        if token not in stop_words:
            result_tokens.append(token)
    return result_tokens
```

```

from nltk.stem import WordNetLemmatizer
from nltk import word_tokenize, pos_tag
nltk.download('averaged_perceptron_tagger')
nltk.download('wordnet')

def lemmatization(token_text):
    """
    This function performs the lemmatization operations as explained above.
    Input Args:
    token_text: list of tokens.
    Returns:
    lemmatized_tokens: list of lemmatized tokens.
    """
    lemma_tokens = []
    wordnet = WordNetLemmatizer()
    lemmatized_tokens = [wordnet.lemmatize(token, pos = 'v') for token in token_t

    return lemmatized_tokens

```

```

➡ [nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data]    /root/nltk_data...
[nltk_data]    Package averaged_perceptron_tagger is already up-to-
[nltk_data]    date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]    Package wordnet is already up-to-date!

```

```

from nltk.stem import PorterStemmer

def stemming(text):
    """
    This function performs stemming operations.
    Input Args:
    token_text: list of tokenize text.
    Returns:
    stemm_tokens: list of stemmed tokens.
    """
    porter = PorterStemmer()
    stemm_tokens = []
    for word in text:
        stemm_tokens.append(porter.stem(word))
    return stemm_tokens

```

```
def lower_order(text):
    """
    This function converts all the text in input text to lower order.
    Input Args:
    token_text : input text.
    Returns:
    small_order_text : text converted to small/lower order.
    """
    small_order_text = text.lower()
    return small_order_text

# Test:
sample_text = "This Is some Normalized TEXT"
sample_small = lower_order(sample_text)
print(sample_small)
```

→ this is some normalized text

✓ Build a Text Cleaning Pipeline

```
def text_cleaning_pipeline(dataset, rule = "lemmatize"):
    """
    This function cleans the dataset
    """
    # Convert the input to small/lower order.
    data = lower_order(dataset)
    # Remove URLs
    data = remove_urls(data)
    # Remove emojis
    data = remove_emoji(data)
    # Remove all other unwanted characters.
    data = removeunwanted_characters(data)

    # Create tokens.
    tokens = data.split()
    # Remove stopwords:
    tokens = remove_stopwords(tokens)
    # Stemming or Lemmatization:
    if rule == "lemmatize":
        tokens = lemmatization(tokens)
    elif rule == "stem":
        tokens = stemming(tokens)
    else:
        print("Pick between lemmatize or stem")

    return " ".join(tokens)
```

✓ Text Classification using Machine Learning Models

✓ Instructions: Trump Tweet Sentiment Classification

1. Load the Dataset

Load the dataset named `"trump_tweet_sentiment_analysis.csv"` using `pandas`. Ensure the dataset contains at least two columns: `"text"` and `"label"`.

2. Text Cleaning and Tokenization

Apply a text preprocessing pipeline to the `"text"` column. This should include:

- Lowercasing the text
- Removing URLs, mentions, punctuation, and special characters
- Removing stopwords
- Tokenization (optional: stemming or lemmatization)
- "Complete the above function"

3. Train-Test Split

Split the cleaned and tokenized dataset into **training** and **testing** sets using `train_test_split` from `sklearn.model_selection`.

4. TF-IDF Vectorization

Import and use the `TfidfVectorizer` from `sklearn.feature_extraction.text` to transform the training and testing texts into numerical feature vectors.

5. Model Training and Evaluation

Import **Logistic Regression** (or any machine learning model of your choice) from `sklearn.linear_model`. Train it on the TF-IDF-embedded training data, then evaluate it using the test set.

- Print the **classification report** using `classification_report` from `sklearn.metrics`.

```
import pandas as pd
import numpy as np
```

1. Load the dataset

```
df = pd.read_csv('/content/drive/MyDrive/2025 - 6CS012 - AI and ML - Student/We
```

```
df.head()
```



	text	Sentiment
0	RT @JohnLeguizamo: #trump not draining swamp b...	0
1	ICYMI: Hackers Rig FM Radio Stations To Play A...	0
2	Trump protests: LGBTQ rally in New York https:...	1
3	"Hi I'm Piers Morgan. David Beckham is awful b...	0
4	RT @GlennFranco68: Tech Firm Suing BuzzFeed fo...	0

```
df.rename(columns={"Sentiment": "label"}, inplace=True)
```

```
df.columns
```



```
Index(['text', 'label'], dtype='object')
```

2. Text Cleaning and Tokenization

```
cleaned_tokens = df["text"].apply(lambda dataset: text_cleaning_pipeline(dataset
```

```
df["clean_text"] = cleaned_tokens
```

```
df['clean_text'][0]
```



```
'rtnot drain swamp taxpayer dollars trip advertise properties'
```

3. Train-Test Split

```
from sklearn.model_selection import train_test_split
```

```
X = df['clean_text']  
y = df['label']
```

```
test_size = 0.2  
random_state = 42
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size,
```

4. TF-IDF Vectorization

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
vectorizer = TfidfVectorizer(max_features=500)
```

```
X_train_vec = vectorizer.fit_transform(X_train)
```

```
X_test_vec = vectorizer.transform(X_test)
```

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression()
```

```
model.fit(X_train_vec, y_train)
```



```
▼ LogisticRegression ⓘ ?  
LogisticRegression()
```

```
from sklearn.metrics import classification_report
```

```
y_pred = model.predict(X_test_vec)
```

```
print(classification_report(y_test, y_pred))
```



	precision	recall	f1-score	support
0	0.81	0.92	0.86	248563
1	0.77	0.56	0.65	121462
accuracy			0.80	370025
macro avg	0.79	0.74	0.75	370025
weighted avg	0.80	0.80	0.79	370025

```
df.shape
```



```
(1850123, 3)
```



```

from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV

pipeline = Pipeline([
    ('tfidf', TfidfVectorizer()),
    ('clf', LogisticRegression())
])

param_grid = {
    'tfidf__max_features': [1000, 3000, 5000, 7000, 10000]
}

grid = GridSearchCV(pipeline, param_grid, cv=3, scoring='accuracy', verbose=1)
grid.fit(X_train, y_train)

print("Best max_features:", grid.best_params_)
print("Best accuracy:", grid.best_score_)

```

⇒ Fitting 3 folds for each of 5 candidates, totalling 15 fits
 /usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:4
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(
 /usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:4
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(
 Best max_features: {'tfidf__max_features': 10000}
 Best accuracy: 0.934375291365842

